



# NoSQL Data Management: Concepts and Systems

**November 8, 2018**

**BUAN 6320 Database Foundations**

# History

- SQL Databases were dominant for decades
  - Persistent storage
  - Standards based
  - Concurrency Control
  - Application Integration
  - ACID
  - Designed to run on a single big machine
- Cloud computing changes that dramatically
  - Cluster of machines
  - Large amount of unreliable machines
  - Distributed System
  - Schema-free unstructured Big Data

# Methods to Run a Database


- Virtual Machine Image
  - Users purchase virtual machine instances to run a database on these
  - Upload and setup own image with database, or use ready-made images with optimized database installations
  - E.g. Oracle Database 11g Enterprise Edition image for Amazon EC2 and for Microsoft Azure.
- Database as a service (DBaaS)
  - Using a database without physically launching a virtual machine instance
  - No configuration or management needed by application owners
  - E.g. Amazon Web Services provide SimpleDB, Amazon Relational Database Service (RDS), DynamoDB
- Managed database hosting
  - Not offered as a service, but hosted and managed by the cloud database vendor
  - E.g. Rackspace offers managed hosting for MySQL
- TOSCA
  - Description of Cloud Services as Topology combined with the database stack
  - Vendor-neutral automatic provisioning and management with OpenTOSCA
  - Policies to define security requirements of the Cloud Service
  - Portable and interoperable definition of data security and compliance aspects

# Which Data Model?

- Relational Databases
  - Standard SQL database available for Cloud Environments as Virtual Machine Image or as a service depending on the vendor
  - Not cloud-ready: Difficult to scale
- NoSQL databases
  - Database which is designed for the cloud
  - Built to serve heavy read/write loads
  - Good ability to scale up and down
  - Applications built based on SQL data model require a complete rewrite
  - E.g. Apache Cassandra, CouchDB and MongoDB

# How to scale the data management?

- Vertical scaling – Scale up



- Horizontal scaling – Scale out



# Who Uses NoSQL?

Google – Big Table, Google Apps,  
Google Search



Facebook – Social network



Twitter



Amazon – DynamoDB and SimpleDB



CERN

GitHub





# Definition and Goals of NoSQL databases

- No formal NoSQL definition available!
- Store very large scale data called “Big data”
- Typically scale horizontally
- Simple query mechanisms
- Often designed and set up for a concrete application
- Typical NoSQL characteristics:
  - Non-relational
  - Schema-free
  - Open Source
  - Simple API
  - Distributed
  - Eventual consistency



# Non-relational are Schema-free

- NoSQL databases generally do not follow the relational model
- Do not provide tables with flat fixed-column records
- Work with self-contained (hierarchical) aggregates or BLOBs
- No need for object-relational mapping and data normalization
- No complex and costly features like query languages, query planners, referential integrity, joins, ACID
- Most NoSQL databases are schema-free or have relaxed schemas
- No need for definition of any sort of schema of the data
- Allows heterogeneous structures of data in the same domain




# Simple API and Distributed


- Often simple interfaces for storage and querying data provided
  - APIs often allow low-level data manipulation and selection methods
  - Often no standard based query language is used
  - Text-based protocols often using HTTP REST with JSON
  - Web-enabled databases running as internet-facing services
- 
- Several NoSQL databases can be executed in a distributed fashion
  - Providing auto-scaling and fail-over capabilities
  - Often ACID is sacrificed for scalability and throughput
  - Often no synchronous replication between distributed nodes is possible, e.g. asynchronous Multi-Master Replication, peer-to-peer, HDFS Replication
  - Only providing eventual consistency

# Core Categories of NoSQL Systems

- **Key-Value Stores**  
Manage associative arrays  
Big hash table
- **Wide Column Stores**  
Each storage block contains only data from one column  
Read and write is done using columns  
(rather than rows – like in SQL)
- **Document Stores**  
Store documents consisting of tagged values  
Data is a collection of key value pairs  
Provides structure and encoding of the managed data  
Encoded using XML, JSON, BSON  
Schema-free
- **Graph DB**  
Network database using graphs with node and edges for storage  
Nodes represent entities, edges represent their relationships



Row ID	Columns...		
1	Name	Website	
	Preston	www.example.com	
2	Name	Website	
	Julia	www.example.com	
3	Name	Email	Website
	Alice	example@example.com	www.example.com



# SQL and NoSQL Systems

**Relational:** [MySQL](#), [PostgreSQL](#), [SQLite](#), [Firebird](#), [MariaDB](#), [Oracle DB](#), [SQL server](#), [IBM DB2](#), [IBM Informix](#), [Teradata](#)

**Key value-stores:** [Memcachedb](#), [Redis](#), [Riak](#), [Amazon DynamoDB](#), [Voldemort](#), [FoundationDB](#), [leveldb](#), [BangDB](#), [KAI](#), [hamsterdb](#), [Tarantool](#), [Maxtable](#), [HyperDex](#), [Genomu](#)

**Column family:** [Big table](#), [Hbase](#), [hyper table](#), [Cassandra](#), [Apache Accumulo](#)

**Document:** [Mongo DB](#), [Couch DB](#), [Rethink DB](#), [Raven DB](#), [terastore](#), [Jas DB](#), [Raptor DB](#), [djon DB](#), [EJDB](#), [denso DB](#), [Couchbase](#)

**Graph databases:** [AllegroGraph](#), [Neo4j](#), [OrientDB](#), [InfiniteGraph](#), [graphbase](#), [sparkledb](#), [flockdb](#), [BrightstarDB](#)

**Object:** [ZODB](#), [DB4O](#), [Eloquera](#), [Versant](#), [Objectivity DB](#), [VelocityDB](#)

**RDF Stores:** [Apache Jena](#), [Sesame](#)

**Multimodel Databases:** [arangodb](#), [Datomic](#), [Orient DB](#), [FatDB](#), [AlchemyDB](#)

**XML Databases:** [BaseX](#), [Sedna](#), [eXist](#)

**Hierarchical:** [InterSystems Caché](#), [GT.M](#)

# SQL to NoSQL Terminology

SQL Terms/Concepts	MongoDB Terms/Concepts
database	<a href="#">database</a>
table	<a href="#">collection</a>
row	<a href="#">document</a> or <a href="#">BSON</a> document
column	<a href="#">field</a>
index	<a href="#">index</a>
table joins	<a href="#">\$lookup</a> , embedded documents
primary key Specify any unique column or column combination as primary key.	<a href="#">primary key</a> In MongoDB, the primary key is automatically set to the <a href="#">_id</a> field.
aggregation (e.g. group by)	aggregation pipeline See the <a href="#">SQL to Aggregation Mapping Chart</a> .
transactions	<a href="#">transactions</a> <b>TIP</b> For many scenarios, the <a href="#">denormalized data model (embedded documents and arrays)</a> will continue to be optimal for your data and use cases instead of multi-document transactions. That is, for many scenarios, modeling your data appropriately will minimize the need for multi-document transactions.

# NoSQL Overview

- Introduction to NoSQL
- Basic Concepts for NoSQL
  - CAP-Theorem
  - Eventual Consistency
  - Consistent Hashing
  - MVCC-Protocol
  - Query Mechanisms for NoSQL
- Overview of NoSQL Systems

# SQL to NoSQL

The following table presents some database executables and the corresponding MongoDB executables. This table is *not* meant to be exhaustive.

MongoDB	MySQL	Oracle	MS SQL	DB2
Database Server	<a href="#">mongod</a>	mysqld	MSSQL\$SQLEXPRESS	IDS
Database Client	<a href="#">mongo</a>	mysql	SSMS	DB-Access



# GENERAL STRUCTURE OF NOSQL DATA

- {
- \_id:  
    ObjectId("509a8fb2f3f4948bd2f983a0"),
- user\_id: "abc123",
- age: 55,
- status: 'A'
- }

# CREATE TABLE

The following table presents the various SQL statements related to table-level actions and the corresponding MongoDB statements.

SQL Schema Statements	MongoDB Schema Statements
<b>CREATE TABLE</b> people ( id MEDIUMINT <b>NOT NULL</b> AUTO_INCREMENT, user_id Varchar(30), age Number, status char(1), <b>PRIMARY KEY</b> (id) )	<p>Implicitly created on first <a href="#">insertOne()</a> or <a href="#">insertMany()</a> operation. The primary key <code>_id</code> is automatically added if <code>_id</code> field is not specified.</p> <pre>db.people.insertOne( { user_id: "abc123", age: 55, status: "A" } )</pre> <p>However, you can also explicitly create a collection:</p> <pre>db.createCollection("people")</pre>
<b>DROP TABLE</b> people	<pre>db.people.drop()</pre>

# ALTER TABLE ADD COLUMN

## SQL Schema Statements

**ALTER TABLE** people **ADD** join\_date  
DATETIME

## MongoDB Schema Statements

Collections do not describe or enforce the structure of its documents; i.e. there is no structural alteration at the collection level.

However, at the document level, [updateMany\(\)](#) operations can add fields to existing documents using the [\\$set](#) operator.

```
db.people.updateMany(
{ },
{ $set: { join_date: new Date() } }
)
```

# ALTER TABLE DROP COLUMN

## SQL Schema Statements

**ALTER TABLE** people **DROP COLUMN**  
join\_date

## MongoDB Schema Statements

Collections do not describe or enforce the structure of its documents; i.e. there is no structural alteration at the collection level.

However, at the document level, [updateMany\(\)](#) operations can remove fields from documents using the [\\$unset](#) operator.

```
db.people.updateMany(  
  { }, { $unset: { "join_date": "" } }  
)
```

# CREATE INDEX

## SQL Schema Statements

```
CREATE INDEX idx_user_id_asc ON  
people(user_id)
```

```
CREATE INDEX idx_user_id_asc_age_desc  
ON people(user_id, age DESC)
```

```
DROP TABLE people
```

## MongoDB Schema Statements

```
db.people.createIndex( { user_id: 1 } )
```

```
db.people.createIndex( { user_id: 1, age: -1 } )
```

```
db.people.drop()
```

# INSERT SQL AND NoSQL

The following table presents the various SQL statements related to inserting records into tables and the corresponding MongoDB statements.

## SQL Schema Statements

```
INSERT INTO people(user_id, age, status)
VALUES ("bcd001", 45, "A")
```

## MongoDB Schema Statements

```
db.people.insertOne(
{ user_id: "bcd001", age: 45, status: "A" }
)
```



# SELECTS

The following table presents the various SQL statements related to reading records from tables and the corresponding MongoDB statements.

**NOTE:** The [find\(\)](#) method always includes the `_id` field in the returned documents unless specifically excluded through [projection](#). Some of the SQL queries below may include an `_id` field to reflect this, even if the field is not included in the corresponding [find\(\)](#) query.

SQL SELECT Statements	MongoDB find() Statements
<b>SELECT</b> * <b>FROM</b> people	db.people.find()
<b>SELECT</b> id, user_id, status <b>FROM</b> people	db.people.find( { }, { user_id: 1, status: 1 } )

# SELECTS, Cont.

SQL Schema Statements	MongoDB Schema Statements
<b>SELECT</b> user_id, status <b>FROM</b> people	db.people.find( { }, { user_id: 1, status: 1, _id: 0 } )
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> status = "A"	db.people.find( { status: "A" } )
<b>SELECT</b> user_id, status <b>FROM</b> people <b>WHERE</b> status = "A"	db.people.find( { status: "A" }, { user_id: 1, status: 1, _id: 0 } )

# SELECTS, Cont.

SQL Schema Statements	MongoDB Schema Statements
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> status != "A"	db.people.find( { status: { \$ne: "A" } } )
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> status = "A" <b>AND</b> age = 50	db.people.find( { status: "A", age: 50 } )
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> status = "A" <b>OR</b> age = 50	db.people.find( { \$or: [ { status: "A" } , { age: 50 } ] } )
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> age > 25	db.people.find( { age: { \$gt: 25 } } )

# SELECTS, Cont.

## SQL Schema Statements

**SELECT** \* **FROM** people **WHERE** age < 25

**SELECT** \* **FROM** people **WHERE** age > 25  
**AND** age <= 50

**SELECT** \* **FROM** people **WHERE** user\_id **like**  
"%bc%"

## MongoDB Schema Statements

db.people.find( { age: { \$lt: 25 } } )

db.people.find( { age: { \$gt: 25, \$lte: 50 } } )

db.people.find( { user\_id: /bc/ } )

-or-

db.people.find( { user\_id: { \$regex: /bc/ } } )

# SELECTS, Cont.

SQL Schema Statements	MongoDB Schema Statements
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> user_id <b>like</b> "bc%"	db.people.find( { user_id: /^bc/ } ) -or- db.people.find( { user_id: { \$regex: /^bc/ } } )
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> status = "A" <b>ORDER BY</b> user_id <b>ASC</b>	db.people.find( { status: "A" } ).sort( { user_id: 1 } )
<b>SELECT</b> * <b>FROM</b> people <b>WHERE</b> status = "A" <b>ORDER BY</b> user_id <b>DESC</b>	db.people.find( { status: "A" } ).sort( { user_id: -1 } )

# SELECTS, Cont.

SQL Schema Statements	MongoDB Schema Statements
<b>SELECT COUNT(*) FROM</b> people	db.people.count()  <i>or</i>  db.people.find().count()
<b>SELECT COUNT</b> (user_id) <b>FROM</b> people	db.people.count( { user_id: { \$exists: <b>true</b> } } )  <i>or</i>  db.people.find( { user_id: { \$exists: <b>true</b> } } ).count()



# SELECTS, Cont.

## SQL Schema Statements

**SELECT COUNT(\*) FROM** people **WHERE** age  
> 30

**SELECT DISTINCT**(status) **FROM** people

## MongoDB Schema Statements

db.people.count( { age: { \$gt: 30 } } )

or

db.people.find( { age: { \$gt: 30 } } ).count()

db.people.aggregate( [ { \$group : { \_id :  
"\$status" } } ] )

or, for distinct value sets that do not exceed  
the [BSON](#)  
[size limit](#)

db.people.distinct( "status" )

# SELECTS, Cont.

SQL Schema Statements	MongoDB Schema Statements
<b>SELECT</b> * <b>FROM</b> people <b>LIMIT</b> 1	db.people.findOne()  or  db.people.find().limit(1)
<b>SELECT</b> * <b>FROM</b> people <b>LIMIT</b> 5 <b>SKIP</b> 10	db.people.find().limit(5).skip(10)
<b>EXPLAIN SELECT</b> * <b>FROM</b> people <b>WHERE</b> status = "A"	db.people.find( { status: "A" } ).explain()

# Update Records

The following table presents the various SQL statements related to updating existing records in tables and the corresponding MongoDB statements.

SQL Update Statements	MongoDB updateMany() Statements
<b>UPDATE</b> people <b>SET</b> status = "C" <b>WHERE</b> age > 25	db.people.updateMany( { age: { \$gt: 25 } }, { \$set: { status: "C" } } )
<b>UPDATE</b> people <b>SET</b> age = age + 3 <b>WHERE</b> status = "A"	db.people.updateMany( { status: "A" } , { \$inc: { age: 3 } } )

# Delete Records

The following table presents the various SQL statements related to updating existing records in tables and the corresponding MongoDB statements.

SQL Delete Statements	MongoDB deleteMany() Statements
<b>DELETE FROM</b> people <b>WHERE</b> status = "D"	db.people.deleteMany( { status: "D" } )
<b>DELETE FROM</b> people	db.people.deleteMany({})